

Classification Assignment Report

- **Identify your problem statement:**

Classification assignment report.

- **Tell basic info about the dataset**

Rows: 399

Columns:28

Output variable: Classification

- **Mention the pre-processing method if you're doing any (like converting string to number – nominal data)**

Pre-processing method is Nominal data (one hot encoding) because input contains text, we are using get_dummies method and the parameters we are passing dataset, dtype=int, drop_first=true

- **All the research values of each algorithm should be documented. (You can make tabulation or screenshot of the results.)**

The screenshot shows a Jupyter Notebook titled "1-Logistic-Regression" with a last checkpoint of 24 minutes ago. The interface includes a menu bar (File, Edit, View, Run, Kernel, Settings, Help) and a toolbar with icons for file operations and execution. The main content area displays a confusion matrix and two code cells with their outputs.

	precision	recall	f1-score	support
0	0.96	0.99	0.98	150
1	1.00	0.98	0.99	249
accuracy			0.98	399
macro avg	0.98	0.98	0.98	399
weighted avg	0.98	0.98	0.98	399

```
[240]: from sklearn.metrics import roc_auc_score
       roc_auc_score(dependent, grid.predict_proba(independent)[:,:1])
```

```
[240]: 0.99978580999062918
```

```
[241]: from sklearn.metrics import f1_score
       f1_macro=f1_score(dependent,grid_pred,average='weighted')
       print("The f1_macro value for best parameter {}".format(grid.best_params_),f1_macro)
```

The f1_macro value for best parameter {'penalty': 'l2', 'solver': 'lbfgs'}: 0.9825109776736257

- **Mention your final model, justify why u have chosen the same**

Final model is logistic-regression because the "roc_auc_score" 0.99978580999062918, and accuracy is 98%

Complete report for all algorithms:

Jupyter

1-Logistic-Regression

Last Checkpoint: 24 minutes ago

File

Edit

View

Run

Kernel

Settings

Help

+

✂

▶

■

↺

▶▶

Code

▼

```
[[149  1]
 [ 6 243]]
```

	precision	recall	f1-score	support
0	0.96	0.99	0.98	150
1	1.00	0.98	0.99	249
accuracy			0.98	399
macro avg	0.98	0.98	0.98	399
weighted avg	0.98	0.98	0.98	399

```
[240]: from sklearn.metrics import roc_auc_score
roc_auc_score(dependent, grid.predict_proba(independent)[:,:1])

[240]: 0.9997858099062918

[241]: from sklearn.metrics import f1_score
f1_macro=f1_score(dependent,grid_pred,average='weighted')
print("The f1_macro value for best parameter {}".format(grid.best_params_),f1_macro)

The f1_macro value for best parameter {'penalty': 'l2', 'solver': 'lbfgs'}: 0.9825109776736257
```

WhatsApp

Home

2-SVM

2236907-Classification_Ass

localhost:8888/notebooks/Week4.1-Machine-Learning-Classification/3.Assignment/2-SVM.ipynb?

Jupyter

2-SVM

Last Checkpoint: 14 hours ago

File

Edit

View

Run

Kernel

Settings

Help

+

✂

▶

■

↺

▶▶

Code

▼

```
[[149  1]
 [ 9 240]]
```

	precision	recall	f1-score	support
0	0.94	0.99	0.97	150
1	1.00	0.96	0.98	249
accuracy			0.97	399
macro avg	0.97	0.98	0.97	399
weighted avg	0.98	0.97	0.98	399

```
[6]: from sklearn.metrics import roc_auc_score
roc_auc_score(dependent, grid.predict_proba(independent)[:,:1])

[6]: 0.999330655957162

[7]: from sklearn.metrics import f1_score
f1_macro=f1_score(dependent,grid_pred,average='weighted')
print("The f1_macro value for best parameter {}".format(grid.best_params_),f1_macro)

The f1_macro value for best parameter {'kernel': 'rbf'}: 0.9750582392902479
```

Jupyter 3.DecisionTreeClassification Last Checkpoint: 14 hours ago

File Edit View Run Kernel Settings Help

Code

```
print(cf_report)
[[150  0]
 [ 0 249]]
      precision    recall  f1-score   support

      0       1.00      1.00      1.00        150
      1       1.00      1.00      1.00        249

   accuracy               1.00           399
  macro avg       1.00      1.00      1.00           399
 weighted avg       1.00      1.00      1.00           399
```

```
[7]: from sklearn.metrics import roc_auc_score
      roc_auc_score(dependent, grid.predict_proba(independent)[:,:1])
```

[7]: 1.0

```
[8]: from sklearn.metrics import f1_score
      f1_macro=f1_score(dependent,y_pred,average='weighted')
      print("The f1_macro value for best parameter {}".format(grid.best_params_),f1_macro)

      The f1_macro value for best parameter {'criterion': 'entropy', 'splitter': 'random'}: 1.0
```

Jupyter 4.RandomForest Last Checkpoint: 14 hours ago

File Edit View Run Kernel Settings Help

Code

```
[[150  0]
 [ 0 249]]
      precision    recall  f1-score   support

      0       1.00      1.00      1.00        150
      1       1.00      1.00      1.00        249

   accuracy               1.00           399
  macro avg       1.00      1.00      1.00           399
 weighted avg       1.00      1.00      1.00           399
```

```
[7]: from sklearn.metrics import roc_auc_score
      roc_auc_score(dependent, grid.predict_proba(independent)[:,:1])
```

[7]: 1.0

```
[8]: from sklearn.metrics import f1_score
      f1_macro=f1_score(dependent,grid_pred, average='weighted')
      print("The f1_macro value for best parameter {}".format(grid.best_params_),f1_macro)

      The f1_macro value for best parameter {'criterion': 'log_loss', 'max_features': 'log2'}: 1.0
```

Jupyter 5.KNN Last Checkpoint: 14 hours ago

File Edit View Run Kernel Settings Help

Code

```
[26]: print(cf_metrics)
      print(cf_report)
```

```
[[149  1]
 [  5 244]]
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	150
1	1.00	0.98	0.99	249
accuracy			0.98	399
macro avg	0.98	0.99	0.98	399
weighted avg	0.99	0.98	0.99	399

```
[27]: from sklearn.metrics import roc_auc_score
      roc_auc_score(dependent, grid.predict_proba(independent)[:,:1])
```

```
[27]: 0.999437751004016
```

```
[34]: from sklearn.metrics import f1_score
      f1_macro=f1_score(dependent,grid_pred, average='weighted')
      print("The f1_macro value for best parameter {}".format(grid.best_params_),f1_macro)
```

The f1_macro value for best parameter {'algorithm': 'auto', 'weights': 'uniform'}: 0.9850004566071048

Jupyter 6.Bernaulli's NB Last Checkpoint: 14 hours ago

File Edit View Run Kernel Settings Help

Code

```
[6]: print(cf_matrix)
```

```
[[149  1]
 [  9 240]]
```

```
[7]: print(cf_report)
```

	precision	recall	f1-score	support
0	0.94	0.99	0.97	150
1	1.00	0.96	0.98	249
accuracy			0.97	399
macro avg	0.97	0.98	0.97	399
weighted avg	0.98	0.97	0.98	399

```
[8]: from sklearn.metrics import roc_auc_score
      roc_auc_score(dependent, grid.predict_proba(independent)[:,:1])
```

```
[8]: 0.991285140562249
```

```
[9]: from sklearn.metrics import f1_score
      f1_macro=f1_score(dependent,grid_pred, average='weighted')
      print("The f1_macro value for best parameter {}".format(grid.best_params_),f1_macro)
```

The f1_macro value for best parameter {'alpha': 0.1, 'binarize': 0.0}: 0.9750582392902479

localhost:8888/notebooks/Week4.1-Machine-Learning-Classification/3.Assignment/7.CategoricalNB.ipynb

Jupyter 7.CategoricalNB Last Checkpoint: 13 hours ago

File Edit View Run Kernel Settings Help

JupyterLab Python 3 (ipykernel)

```
[21]: print(cf_matrix)
      print(cf_report)

      [[150  0]
       [ 1 248]]

           precision    recall  f1-score   support

              0       0.99      1.00      1.00        150
              1       1.00      1.00      1.00        249

   accuracy          1.00
  macro avg          1.00
 weighted avg          1.00
```

```
[23]: from sklearn.metrics import roc_auc_score
      roc_auc_score(dependent, grid.predict_proba(independent)[:,1])

[23]: 0.9999196787148594

[25]: from sklearn.metrics import f1_score
      f1_macro = f1_score(dependent, grid_pred, average = 'weighted')
      print("The f1_macro value for best parameter {}: {}".format(grid.best_params_,f1_macro)

      The f1_macro value for best parameter {'alpha': 1.0, 'class_prior': None, 'fit_prior': True, 'force_alpha': True, 'min_categories': None}: 0.9974953761738116
```

localhost:8888/notebooks/Week4.1-Machine-Learning-Classification/3.Assignment/8.ComplementNB.ipynb

Jupyter 8.ComplementNB Last Checkpoint: 14 hours ago

File Edit View Run Kernel Settings Help

JupyterLab Python 3 (ipykernel)

```
[19]: print(cf_matrix)
      print(cf_report)

      [[145  5]
       [ 39 210]]

           precision    recall  f1-score   support

              0       0.79      0.97      0.87        150
              1       0.98      0.84      0.91        249

   accuracy          0.89
  macro avg          0.88
 weighted avg          0.91
```

```
[20]: from sklearn.metrics import roc_auc_score
      roc_auc_score(dependent, grid.predict_proba(independent)[:,1])

[20]: 0.9623025435073628

[21]: from sklearn.metrics import f1_score
      f1_score = f1_score(dependent, grid_pred, average = 'weighted')
      print("evaluation metrics {}: {}".format(grid.best_params_,f1_score)

      evaluation metrics {'alpha': 0.1, 'fit_prior': True, 'force_alpha': True}: 0.8912968721618211
```

What x Home x 9. GaussianNB x 8. Correlation x 7. Categorical x 6. Bernoulli x 5. K-Nearest x 4. Random Forest

localhost:8888/notebooks/Week4.1-Machine-Learning-Classification/3.Assignment/9.%20GaussianNB.ipynb

Jupyter 9. GaussianNB Last Checkpoint: 14 hours ago

File Edit View Run Kernel Settings Help

+ × Copy Paste Run Cell Code

```
[26]: print(cf_matrix)
      print(cf_report)
```

```
[[ 35 115]
 [ 18 231]]
```

	precision	recall	f1-score	support
0	0.66	0.23	0.34	150
1	0.67	0.93	0.78	249
accuracy			0.67	399
macro avg	0.66	0.58	0.56	399
weighted avg	0.66	0.67	0.61	399

```
[31]: from sklearn.metrics import roc_auc_score
      roc_auc_score(dependent, grid.predict_proba(independent)[:,:1])
```

```
[31]: 0.6322623828647925
```

```
[33]: from sklearn.metrics import f1_score
      f1_score = f1_score(dependent, grid_pred, average = 'weighted')
      print("evaluation metrics {}".format(grid.best_params_),f1_score)
```

```
evaluation metrics {'priors': None, 'var_smoothing': 0.1}: 0.6141987829614605
```

What x Home x 10. MultinomialNB x 9. GaussianNB x 8. Correlation x 7. Categorical x 6. Bernoulli x 5. K-Nearest x 4. Random Forest

localhost:8888/notebooks/Week4.1-Machine-Learning-Classification/3.Assignment/10.MultinomialNB.ipynb

Jupyter 10. MultinomialNB Last Checkpoint: 13 hours ago

File Edit View Run Kernel Settings Help

+ × Copy Paste Run Cell Code

```
[20]: print(cf_matrix)
      print(cf_report)
```

```
[[145   5]
 [ 54 195]]
```

	precision	recall	f1-score	support
0	0.73	0.97	0.83	150
1	0.97	0.78	0.87	249
accuracy			0.85	399
macro avg	0.85	0.87	0.85	399
weighted avg	0.88	0.85	0.85	399

```
[23]: from sklearn.metrics import roc_auc_score
      roc_auc_score(dependent, grid.predict_proba(independent)[:,:1])
```

```
[23]: 0.9499866131191431
```

```
[25]: from sklearn.metrics import f1_score
      f1_score = f1_score(dependent, grid_pred, average = 'weighted')
      print("evaluation metrics {}".format(grid.best_params_),f1_score)
```

```
evaluation metrics {'alpha': 1.0, 'force_alpha': True}: 0.8544422491701906
```